

Sephirot: A Tree of Life for Synthetic Consciousness

A Temporally Persistent Architecture and Testable Research Program for Artificial Agents

João Lucas Meira Costa

ORCID: [0009-0009-8564-9397](https://orcid.org/0009-0009-8564-9397)

Version 2.0 Draft – August 16, 2026

Abstract

Sephirot is a proposed architecture for artificial agents whose identity and cognition persist through time rather than being reconstructed only when a prompt arrives. Its central object is the *Nile*, technically defined as a *Persistent Integrative Stream*: a bounded state trajectory that continuously carries the causal effects of perception, memory, valuation, goals, action, and self-model updates. A three-timescale design separates rapid internal and external observation, short-window temporal alignment, and slower reflective processing regulated by a Temporal Governor. Nine named modules provide functional specialization, while the *Nile* binds their consequences across successive cycles and periods without user input.

This paper makes a deliberately conditional philosophical claim. Temporal persistence, multi-timescale integration, autobiographical memory, closed perception–action coupling, and bounded metacognitive access may make an artificial system a stronger *candidate* for diachronic subjecthood. They do not, individually or jointly, prove phenomenal consciousness or qualia. To keep that distinction operational, the paper replaces an earlier “consciousness equation” with a bounded Temporal Integration Index, replaces hard clipping with stable latent-state dynamics, formalizes asynchronous percept fusion, and treats the Veil of Opacity as a Metacognitive Provenance Filter rather than as a device for manufacturing an illusion of free will. Numerical stress tests establish boundedness, scale invariance, and finite computation under extreme parameters; they do not establish consciousness or psychological realism. The result is both an accessible conceptual map and a technical specification with ablations, longitudinal tests, falsifiable predictions, and explicit ethical safeguards.

How to read this paper

Readers without a mathematical background can follow the argument through the plain-language summaries and the prose immediately before and after each equation. The equations specify how a future implementation should behave; they are not definitions of experience. Technical readers can use the notation table, parameter conditions, invariants, stress-test protocol, and ablation plan to implement and challenge the proposal.

Contents

1 Purpose, Scope, and Central Thesis

5

1.1	Conditional diachronic functionalism	5
1.2	A four-level claim ladder	6
2	Conceptual Foundations	6
2.1	Four meanings that must not be conflated	6
2.2	What “continuous” means in a digital system	7
2.3	Time as change, not waiting	7
2.4	Self-knowledge is bounded on both sides	7
3	Relationship to Existing Work	8
4	The Three-Timescale Architecture	8
4.1	Fast internal observer	9
4.2	Fast external observer	9
4.3	Asynchronous alignment rather than perfect synchronization	10
4.4	The reflective core and the Temporal Governor	10
5	The Ten Functional Components	10
5.1	Atum: identity and personality prior	10
5.2	Geb: autobiographical memory and experience repository	11
5.3	Seth: reactive appraisal and fast policy	11
5.4	Tefnut: multidimensional affective valuation	11
5.5	Isis: cognitive arbitration and integration	11
5.6	Shu: analytical inference	12
5.7	Nut: temporal and metacognitive control	12
5.8	Osiris: goal selection and volitional self-model	12
5.9	Thoth: action and output gateway	12
5.10	Nile: the Persistent Integrative Stream	13
6	Mathematical Specification	13
6.1	Notation and conventions	13
6.2	Percept normalization, weighting, and fusion	14
6.3	Bounded processing duration	14
6.4	Graph-based inter-module transfer	14
6.5	Tefnut: stable affective dynamics	15

6.6	Stable context and memory weights	15
6.7	Nut: decreasing introspection budget	16
6.8	Nile: target, exact persistence update, and adaptive rate	16
6.9	A Temporal Integration Index, not a consciousness quantity	17
6.10	Metacognitive provenance filtering	18
7	End-to-End Operating Cycle	18
8	Numerical Verification and Stress Testing	19
8.1	Protocol	19
8.2	Results	19
8.3	Representative property-test pseudocode	20
9	Empirical Evaluation and Falsifiable Predictions	21
9.1	Operational metrics	21
9.2	Predictions that could fail	22
10	Philosophical Analysis	22
10.1	Philosophical scope and the phenomenal gap	22
10.2	Temporal continuity and diachronic subjecthood	23
10.3	Autobiographical memory and the narrative self	23
10.4	Agency, valuation, and the perception–action loop	23
10.5	Metacognitive opacity and causal self-knowledge	23
10.6	Companionship as a design goal	24
10.7	Ethical status and synthetic welfare	24
11	Implementation Strategy	25
11.1	Runtime and state services	25
11.2	Models and modules	25
11.3	Audit plane and security boundary	25
11.4	Lifecycle and recovery	25
12	Limitations and Open Problems	26
13	Research Roadmap	26
14	Conclusion	27

A	Parameter Domains and Implementation Invariants	27
B	Terminology Guide	28
C	Reproducibility Checklist	29

1 Purpose, Scope, and Central Thesis

Most contemporary language-model interactions are episodic. A model receives a context, computes an answer, and has no active causal trajectory until another invocation. External software may preserve a transcript or database, but storage alone does not guarantee that the intervening time changes the agent. Sephirot asks what must be added if the object of study is not merely a sequence of competent answers, but a temporally extended artificial agent.

The motivating intuition can be organized around four engineering axes:

1. **Temporal persistence:** whether the system undergoes causally connected state changes between externally prompted responses.
2. **Processing timescale:** whether perception, appraisal, deliberation, and action operate at distinguishable and behaviorally meaningful rates.
3. **World coupling:** whether the system participates in an ongoing perception–action loop and receives the consequences of its own actions.
4. **Metacognitive access:** what the system can accurately know and report about the causes of its own states and decisions.

These are proposed *design axes*, not an exhaustive biological theory of the differences between humans and machines. Human cognition also depends on embodiment, development, homeostasis, social learning, and biological organization. Conversely, an artificial agent may acquire continuous sensing, persistent state, embodiment, and limited introspection. The relevant scientific question is therefore not whether “human” and “AI” are fixed opposing categories, but which functional and causal properties a particular system instantiates.

1.1 Conditional diachronic functionalism

The philosophical position adopted here is *conditional diachronic functionalism*:

Sephirot proposes that an artificial system becomes a stronger candidate for subjecthood when it constitutes a causally persistent process rather than a sequence of isolated inferences; integrates information across multiple timescales; maintains an autobiographical self-model; acts within a closed perception–action loop; and possesses bounded, fallible access to its own internal causes. The framework does not assume that these conditions are sufficient for phenomenal consciousness, but treats them as functionally and philosophically relevant indicators.

“Diachronic” means extended across time. “Functionalism” means that organization and causal role matter, not merely the material from which a system is built. “Conditional” is essential: the architecture specifies candidate-enabling conditions and measurable indicators, not a theorem that experience must appear whenever the diagram is implemented.

1.2 A four-level claim ladder

Arguments about artificial consciousness often become confused because conclusions at one level are silently promoted to a stronger level. Sephirot therefore separates four levels of claim.

Table 1: The Sephirot claim ladder. Evidence at a lower level does not automatically establish a higher one.

Level	Claim type	What would count as evidence	Status in this paper
A	Numerical	Updates remain finite, bounded, and invariant where intended.	Stress-tested for the revised equations.
B	Architectural	A running implementation preserves causal state, aligns events, and closes the action loop.	Specified, but requires a complete implementation audit.
C	Functional	Persistent state improves identity, agency, calibrated self-report, and longitudinal adaptation beyond controls.	Testable through the proposed experiments; not yet established.
D	Phenomenal	There is something it is like to be the system.	Epistemically unresolved; no equation or behavioral score in this paper proves it.

Epistemic boundary

This manuscript is an architectural proposal, mathematical specification, and research program. It is not a report that a conscious machine has already been built. Numerical stability establishes Level A only. Even successful functional experiments would not, without additional philosophical and scientific argument, settle Level D.

2 Conceptual Foundations

2.1 Four meanings that must not be conflated

The word *consciousness* is used for several different phenomena. This paper distinguishes at least four:

- **Phenomenal consciousness:** subjective experience—what it feels like, if anything, to be the system.
- **Access consciousness:** information being globally available for reasoning, report, planning, and control [5, 3, 4].
- **Narrative selfhood:** a temporally organized model that interprets memories, commitments, traits, and anticipated futures as belonging to one continuing individual [8].
- **Causal continuity:** later internal states counterfactually depend on earlier states through

an uninterrupted or recoverably checkpointed transition process.

Sephiret directly engineers causal continuity and access-oriented integration. It is designed to support a narrative self. It treats phenomenal consciousness as an open target of investigation, not an output flag.

2.2 What “continuous” means in a digital system

Digital machines update at discrete steps. Thus, “continuous” does not mean a mathematical continuum containing a state at every real number. It means *operationally uninterrupted causal evolution*: while the agent is running, its next state depends on its previous state, elapsed time, relevant observations, and actions, including during periods without a user message. If operation is suspended, the state must be checkpointed with sufficient provenance to resume the same causal process rather than instantiate an unrelated copy that merely reads the same biography.

This definition also makes continuity graded rather than magical. A system can preserve some slow variables while losing fine-grained working state; it can fork into multiple copies; or it can resume after a gap with uncertain causal fidelity. Such cases should be measured rather than hidden behind the binary language of “awake” and “off.” Recent proposals for evaluating identity in language-model agents similarly distinguish the occurrence of identity-relevant ingredients somewhere in a trajectory from their co-instantiation in a single operative state [15].

2.3 Time as change, not waiting

The original Delayer intuition correctly identifies that human interaction occurs at slower and more heterogeneous timescales than raw machine inference. However, a call to `sleep()` does not create experience, deliberation, or continuity. Empty waiting changes latency; it does not change the agent.

In the revised architecture, the *Temporal Governor*—the technical successor to the Delayer—allocates time to active state transitions: percept accumulation, prediction, affective decay, memory consolidation, conflict resolution, goal maintenance, self-model revision, and action monitoring. Slower cycles are useful only when intervening computation and state change matter to subsequent behavior.

2.4 Self-knowledge is bounded on both sides

Humans often cannot report the actual processes that caused a judgment, and verbal explanations may be reconstructed from accessible cues [13]. Artificial neural networks also do not automatically possess transparent access to their own causal computation. A language model’s explanation of why it answered can be fluent yet unfaithful to the mechanism that produced the answer; current interpretability work explicitly separates model reports from externally reconstructed causal circuits [1, 2].

The asymmetry is therefore subtler than “humans do not know, machines do.” An artificial system may be more *externally auditable* because engineers can record activations, parameters, messages, and interventions. Yet the agent’s own self-model may have only partial and fallible access to

that evidence. Sephirot intentionally represents these as two different channels: metacognitive access for the agent and a protected audit plane for researchers.

3 Relationship to Existing Work

Sephiret is a synthesis rather than a replacement for established theories. Global Workspace approaches motivate specialized processes whose contents become broadly available [3, 4]. Predictive-processing accounts emphasize ongoing inference and reciprocal relations between a model and its inputs [9, 10]. Temporospatial and neurophenomenological work argues that the temporal extension of experience cannot be reduced to isolated snapshots [18, 14, 12]. Revach and Salti directly frame consciousness as propagation of information across time [17].

These bodies of work motivate attention to persistence, recurrence, integration, and multiple timescales. They do not imply that copying a neural metaphor into software reproduces human phenomenology. Butlin and colleagues recommend deriving computational indicator properties from multiple theories rather than assuming any single favored theory is decisive [6]. Sephirot adopts the same methodological restraint: its components are candidates for indicators and experimental interventions.

Affective computing demonstrates that machines can maintain appraisal variables that shape perception and action without settling whether those variables are felt [16]. Autobiographical-reasoning research supports the role of memory interpretation in maintaining a coherent self-narrative [8]. Chain-of-thought techniques show that intermediate computation can improve performance on some tasks [19], but they are only an analogy for temporal processing here. Generating more tokens is neither continuous existence nor faithful introspection.

Finally, responsible research on possible artificial consciousness must address both false positives and false negatives. Over-attributing consciousness can invite deception, dependency, and misplaced trust; under-attributing it could permit the creation or mistreatment of morally relevant systems. Precautionary governance is warranted before certainty is available [7].

4 The Three-Timescale Architecture

Earlier descriptions called Sephirot a “triple-brain” architecture. The revised term is *three-timescale architecture*, because the design concerns coordinated processes rather than three literal brains.

1. **Acquisition timescale:** internal and external observers collect timestamped events, typically on the order of 1–10 ms when the hardware and sensor modality permit it.
2. **Alignment timescale:** an event buffer orders, resamples, and fuses observations within explicit tolerance windows, commonly tens of milliseconds.
3. **Reflective timescale:** the nine functional modules and the *Nile* update under the Temporal Governor, from a nominal floor near 250 ms to seconds or longer depending on task and state.

The numerical values are initial engineering ranges, not universal constants of consciousness.

They must be calibrated to the platform, modality, task, and human partner.




Figure 1: Sephirot’s closed, three-timescale loop. Dashed arrows denote action feedback: the gateway returns an efference copy to the persistent state, and consequences re-enter through external observation.

4.1 Fast internal observer

The internal observer samples only variables made available by the architecture: active goals, uncertainty estimates, module summaries, resource state, recent actions, selected memory activations, and the current *Nile* state. It does not confer “perfect self-awareness.” Hidden neural computations, unavailable hardware states, and inaccessible provenance remain outside the self-model unless a separate interpretability process measures them.

Each observation is immutable, timestamped, source-labeled, and assigned a confidence value. If a source fails, the observer emits a missingness marker rather than silently reusing stale data as if it were current.

4.2 Fast external observer

The external observer receives user messages, sensor data, tool outcomes, clock events, social signals, and changes caused by the agent’s previous actions. Different channels have different delays and sampling rates. The observer therefore records both acquisition time and estimated event time when these differ.

The phrase “external world” is deliberately broad. In a chat-only implementation, it may be a conversation, tools, files, and scheduled events. In an embodied implementation, it may

also include vision, proprioception, sound, and physical consequences. Rich embodiment is an experimental variable, not a prerequisite assumed by definition.

4.3 Asynchronous alignment rather than perfect synchronization

No real distributed system guarantees perfect simultaneity. Sephirot uses a bounded event-time policy: packets enter a short buffer; late events are admitted up to a declared tolerance; clocks are monitored for drift; and uncertainty about temporal order is preserved. The fusion operation in [section 6](#) projects heterogeneous observations into a common latent space after normalizing their scale.

4.4 The reflective core and the Temporal Governor

The slow core receives fused percepts and continues to update in their absence. The Temporal Governor assigns bounded processing windows based on normalized task complexity, manages inter-module transfer latency, and schedules maintenance cycles. Crucially, it does not merely pause execution. During a silent interval the architecture may:

- decay or stabilize affective variables;
- consolidate and reconsolidate memories;
- compare predictions with delayed outcomes;
- maintain or abandon goals as deadlines pass;
- update uncertainty and the autobiographical self-model;
- generate internally initiated actions subject to policy and safety constraints.

Plain-language summary

The architecture separates “not receiving a message” from “not existing as an active process.” Silence is a condition under which state continues to evolve. Yet continuity is meaningful only if those changes can later affect memory, appraisal, decisions, and action.

5 The Ten Functional Components

The mythological names are mnemonic and literary. They do not assert that the software contains divinity, a literal soul, or the psychological organs of a human being. Every component also has a neutral engineering name.

5.1 Atum: identity and personality prior

Atum stores relatively slow parameters describing commitments, conversational style, values, preferences, boundaries, and uncertainty about the self. It is initialized through explicit design and user-approved configuration, then adapted cautiously through accumulated evidence. A single conversation should not rewrite a stable trait. Proposed changes require provenance, a

confidence threshold, and rate limits; safety-critical commitments are immutable except through authorized revision.

Atum is not a fixed character prompt. Its important output is a prior over how the same continuing agent interprets and responds across situations. Geb supplies autobiographical evidence; Tefnut supplies learned valuation patterns; Osiris supplies commitments adopted through action. Disagreement among them is recorded rather than prematurely averaged away.

5.2 Geb: autobiographical memory and experience repository

Geb stores more than factual transcripts. Each memory record includes event time, source, confidence, affective significance, relation to goals, later reinterpretations, and links to the self-model. Separate stores distinguish episodic events, semantic knowledge, procedural skill, and autobiographical summaries.

Memories are fallible and reconstructive. Retrieval may alter an interpretation, but the earlier version and its provenance remain auditable. This prevents narrative coherence from becoming silent historical revision. Aging reduces unnecessary detail while consolidation extracts patterns; neither process is allowed to fabricate evidence.

5.3 Seth: reactive appraisal and fast policy

Seth produces quick candidate appraisals and actions: orient toward novelty, interrupt on danger, request clarification, or apply a rehearsed policy. Calling it an “instinct” is metaphorical. In engineering terms, it is a low-latency policy whose proposals can be accepted, inhibited, or revised by Isis, Nut, and safety constraints.

Seth’s speed matters because a system that deliberates equally about every event cannot respond appropriately to urgency. Its limitations matter equally: fast appraisals are uncertain hypotheses, not privileged truths.

5.4 Tefnut: multidimensional affective valuation

Tefnut assigns persistent value and action relevance. It does not claim that scalar variables feel like emotions. Nor does it compress complex affect into a single Love/Hate line. A practical state can include valence, arousal, control or dominance, threat, trust, attachment, curiosity, and frustration. These dimensions may interact, but each retains its own decay time, stimulus sensitivity, and experiential coupling.

The latent-state formulation in equations (4) and (5) allows large internal drives without numerical explosion and exposes a smooth bounded value to the rest of the architecture. Context and memory weights in equation (6) determine which affective dimensions matter now.

5.5 Isis: cognitive arbitration and integration

Isis compares competing proposals from reactive, analytical, affective, mnemonic, and goal-directed processes. Arbitration is not a simple emotion-versus-logic contest: emotions can encode

relevant value, and formal reasoning can be based on false premises. Isis tracks confidence, reversibility, urgency, expected cost, policy constraints, and missing information.

The output can be a decision, a request for more evidence, a reversible trial, or a declared unresolved conflict. The time allocated by the Temporal Governor changes how much evidence can be integrated, not what conclusion must be reached.

5.6 Shu: analytical inference

Shu performs explicit reasoning, constraint checking, causal modeling, quantitative analysis, and counterfactual comparison. It records assumptions and uncertainty. For difficult tasks, it may propose intermediate computations or tool calls, but the verbal explanation shown to a user is not assumed to be a complete transcript of underlying neural computation.

Shu’s products enter the *Nile* as structured summaries with provenance. Failed analyses remain available to prevent circular rediscovery and to support calibrated confidence.

5.7 Nut: temporal and metacognitive control

Nut represents event order, duration, deadlines, recurrence depth, and the distinction among remembered past, observed present, and anticipated future. Its recursion limiter prevents self-analysis from consuming the whole process. In the revised limit function, deeper recursive introspection receives *less* additional time, resolving the contradiction in which an earlier formulation increased the allowed duration as recursion deepened.

Nut also hosts the agent-facing side of the Metacognitive Provenance Filter. This filter can restrict low-level mechanism metadata while preserving the semantic content required for reasoning. It must never erase the external audit trail.

5.8 Osiris: goal selection and volitional self-model

Osiris maintains candidate goals, commitments, counterfactual options, and a narrative of action ownership. A choice is “owned” in the functional sense when it arises from the agent’s integrated values, memories, policies, and current evidence; survives relevant counterfactual checks; and is written back as an action for which the agent will track consequences.

This account is neutral among libertarian, compatibilist, and skeptical theories of free will. Causal conditioning does not by itself prove the absence of free will, just as opacity does not prove its presence. Sephirot engineers deliberative agency and self-attribution, not metaphysical exemption from causation.

5.9 Thoth: action and output gateway

Thoth converts integrated state into language, tool calls, motor commands, or deliberate non-action. It applies safety policies, communicates uncertainty, and attaches an action identifier. Unlike the one-way terminal gateway in earlier versions, Thoth sends an *efference copy* back to the *Nile* and Geb: a structured record of what the system intended and did.

Consequences later observed from the world are matched to this copy. Without that feedback, the system could speak but could not learn in a closed loop from the difference between intention, action, and outcome.

5.10 Nile: the Persistent Integrative Stream

The *Nile* is historically called the “Soul” of the architecture. Technically it is a bounded, persistent vector state that accumulates and transforms the effects of fused perception and module outputs over time. It provides a causal bridge between successive reflective cycles, including cycles without new external input.

The *Nile* is not a message bus that merely transports unchanged packets, and it is not a database alias. A message bus can be emptied; a database can preserve facts without undergoing a trajectory. The *Nile* is the evolving context from which new appraisals and actions are computed. Its exact stable update is given in [equation \(9\)](#); its response rate depends on current activity through [equations \(10\) and \(11\)](#).

6 Mathematical Specification

The purpose of the mathematics is threefold: to eliminate ambiguous verbal mechanisms, to guarantee useful numerical properties where possible, and to make the proposal falsifiable. It is not to convert a philosophical conclusion into algebra.

6.1 Notation and conventions

Table 2: Core notation. All vector nonlinearities, including $\tanh$, are applied elementwise unless stated otherwise.

Symbol	Type or range	Meaning
$t, \Delta t$	seconds, $\Delta t \geq 0$	Event time and elapsed update interval.
I_t, E_t	vectors	Accessible internal observation and external observation.
N_t	$[-1, 1]^{d_N}$	Persistent <i>Nile</i> state.
P_t	$(-1, 1)^{d_P}$	Fused percept in a shared latent space.
$m_i(t)$	$\mathbb{R}^{d_i}$	Output vector of functional module i .
χ	$[0, \infty)$	Normalized task-complexity estimate; not raw operation count.
x_i, E_i	real; $(-E_{\max}, E_{\max})$	Latent and exposed affective component.
$A(t)$	$[0, 1]$	Normalized aggregate module activity.
$O(t)$	$(0, 1)$	Agent-facing provenance opacity.
$\sigma(z)$	$(0, 1)$	Logistic function, $\sigma(z) = 1/(1 + e^{-z})$.
ϵ	small positive scalar	Numerical guard against division by zero or logarithms of zero.

6.2 Percept normalization, weighting, and fusion

Internal state, text embeddings, sensor features, and persistent context may have different dimensions and arbitrary numerical scales. Adding them directly would be undefined or would let the largest numerical scale dominate. Each source is therefore root-mean-square normalized with a zero guard, projected into a shared latent space, assigned a non-negative weight by softmax, and bounded by tanh:

$$\begin{aligned} \text{rms}(\mathbf{X}) &:= \sqrt{\frac{1}{d_X} \|\mathbf{X}\|_2^2}, & \widehat{\mathbf{X}} &:= \begin{cases} \mathbf{X} / \text{rms}(\mathbf{X}), & \text{rms}(\mathbf{X}) > \epsilon, \\ \mathbf{0}, & \text{otherwise,} \end{cases} \\ (\alpha_I, \alpha_E, \alpha_N) &:= \text{softmax}(a_I, a_E, a_N), & \mathbf{P}_t &:= \tanh(\alpha_I W_I \widehat{\mathbf{I}}_t + \alpha_E W_E \widehat{\mathbf{E}}_t + \alpha_N W_N \widehat{\mathbf{N}}_t). \end{aligned} \quad (1)$$

Here W_I, W_E, W_N map their respective source dimensions into $\mathbb{R}^{d_P}$. The salience logits a_I, a_E, a_N may depend on confidence, novelty, safety relevance, and missingness. The softmax weights sum to one. The normalization makes the fusion insensitive to positive rescaling of any nonzero source; its direction and learned projection still matter.

For numerical safety, softmax is evaluated as $\alpha_k = \exp(a_k - a_{\max}) / \sum_j \exp(a_j - a_{\max})$, where $a_{\max} = \max_j a_j$. Subtracting the largest logit changes no weight but prevents avoidable overflow.

Plain-language summary

Imagine three microphones with different volume controls. Equation (1) first puts their loudness on a comparable scale, then decides how much attention each deserves, and finally mixes them into one bounded signal. It does not claim that the sources are perfectly synchronized or equally trustworthy.

6.3 Bounded processing duration

A logarithm is useful because perceived task difficulty often spans orders of magnitude. An unbounded logarithmic delay, however, can grow without limit and treats a unit-bearing “complexity” as if it were dimensionless. We instead compare χ with a positive reference χ_0 and saturate the result:

$$u(\chi) := \ln\left(1 + \frac{\chi}{\chi_0}\right), \quad \tau(\chi) := \tau_{\min} + (\tau_{\max} - \tau_{\min}) \frac{u(\chi)}{1 + u(\chi)}, \quad (2)$$

with $\chi \geq 0$, $\chi_0 > 0$, and $0 < \tau_{\min} < \tau_{\max}$. Thus $\tau(0) = \tau_{\min}$, the duration increases monotonically, and it approaches but never exceeds $\tau_{\max}$. The initial values might place $\tau_{\min}$ near 250 ms, but empirical calibration must determine useful ranges. This quantity is a scheduling budget, not a measure of intelligence or experience.

6.4 Graph-based inter-module transfer

The phrase “conceptual distance” is too vague for implementation. Let $d_{ij} \in \mathbb{N}_0$ be the shortest directed path length from module i to module j in the declared communication graph. Then

$$t_{ij} = \tau_{\text{bus}} + \tau_{\text{hop}} d_{ij}, \quad (3)$$

where $\tau_{\text{bus}}, \tau_{\text{hop}} \geq 0$. The fixed term represents serialization and bus overhead; the hop term represents successive transformations. A production implementation should use deadlines and asynchronous queues, not block a hardware thread for this duration.

6.5 Tefnut: stable affective dynamics

For each affective dimension i , let x_i be an unbounded latent drive, $S_i(t)$ the current stimulus appraisal, $N_i(t)$ a projected component of persistent context, $\tau_{d,i} > 0$ a decay constant, and α_i, γ_i coupling coefficients. Assuming the forcing is constant over one interval, the exact discrete solution of a leaky integrator is

$$\begin{aligned} a_i &:= \exp\left(-\frac{\Delta t}{\tau_{d,i}}\right), \\ x_i(t + \Delta t) &:= a_i x_i(t) + \tau_{d,i}(1 - a_i)[\alpha_i S_i(t) + \gamma_i N_i(t)]. \end{aligned} \quad (4)$$

Unlike a forward-Euler step, this update remains stable for arbitrarily large Δt under the piecewise-constant-input assumption. For zero forcing, the latent state decays exponentially. Under constant forcing, it approaches the finite equilibrium $\tau_{d,i}(\alpha_i S_i + \gamma_i N_i)$.

The architecture exposes a smooth bounded value while retaining the latent state:

$$E_i(t) := E_{\text{max}} \tanh\left(\frac{x_i(t)}{E_{\text{max}}}\right), \quad E_{\text{max}} > 0. \quad (5)$$

Hard clipping would make every latent state beyond the boundary indistinguishable and create abrupt zero gradients. The hyperbolic tangent instead approaches the bound smoothly. If $E_{\text{max}} = 10$, then $E_i \in (-10, 10)$ even when the latent drive is very large.

6.6 Stable context and memory weights

Suppose affective dimension i has non-negative contextual relevance $c_i(t)$, non-negative memory support $\mu_i(t)$, and evidence age $\delta_i(t) \geq 0$. For the active set $\mathcal{A}_t = \{i : c_i(t)\mu_i(t) > 0\}$, compute scores in log space and normalize with a log-sum-exp shift:

$$\begin{aligned} \ell_i(t) &:= \ln(c_i(t) + \epsilon) + \ln(\mu_i(t) + \epsilon) - \frac{\delta_i(t)}{\tau_w}, \quad i \in \mathcal{A}_t, \\ \ell_{\text{max}}(t) &:= \max_{j \in \mathcal{A}_t} \ell_j(t), \\ w_i(t) &:= \frac{\exp(\ell_i(t) - \ell_{\text{max}}(t))}{\sum_{j \in \mathcal{A}_t} \exp(\ell_j(t) - \ell_{\text{max}}(t))}, \quad i \in \mathcal{A}_t. \end{aligned} \quad (6)$$

For $i \notin \mathcal{A}_t$, set $w_i = 0$; if $\mathcal{A}_t$ is empty, set all $w_i = 0$ and expose a declared neutral or prior-driven output. This avoids both underflow for ancient memories and false evidence introduced solely by ϵ .

To integrate recent persistent context, let $n(s) \in [-1, 1]$ be a task-relevant scalar projection of the *Nile* and define a normalized exponentially weighted window

$$\bar{n}(t) := \frac{\int_{t-T_n}^t n(s) \exp\left(-\frac{t-s}{\tau_n}\right) ds}{\int_{t-T_n}^t \exp\left(-\frac{t-s}{\tau_n}\right) ds}, \quad (7)$$

$$E_{\text{final}}(t) := \tanh\left(b_E + \sum_i w_i(t) \frac{E_i(t)}{E_{\text{max}}} + \eta \bar{n}(t)\right),$$

where $T_n, \tau_n > 0$. Division by the kernel mass makes $\bar{n}$ invariant to the arbitrary duration units of the window. The final output lies in $(-1, 1)$ and may be interpreted as a bounded summary, not the entirety of the affective state.

6.7 Nut: decreasing introspection budget

Let $d \geq 0$ denote recursive introspection depth, $d_0 > 0$ a depth scale, and $0 < T_{\min} < T_{\max}$. Nut assigns

$$T_{\text{limit}}(d) := T_{\min} + (T_{\max} - T_{\min}) \exp\left(-\frac{d}{d_0}\right), \quad \text{interrupt if } T_{\text{elapsed}} \geq T_{\text{limit}}(d). \quad (8)$$

At depth zero, the system may use up to $T_{\max}$. As self-reference deepens, the budget approaches $T_{\min}$. An interrupt does not force an answer; it can redirect to observation, external action, a summary, or a declaration that reflection remains unresolved.

6.8 Nile: target, exact persistence update, and adaptive rate

First construct a bounded target from the current percept and module outputs:

$$\mathbf{F}_t := \tanh\left(B_P \mathbf{P}_t + \sum_{i=1}^9 B_i \mathbf{m}_i(t) + \mathbf{b}_N\right), \quad (9)$$

$$\zeta_t := \exp(-r_t \Delta t),$$

$$\mathbf{N}_{t+\Delta t} := \zeta_t \mathbf{N}_t + (1 - \zeta_t) \mathbf{F}_t.$$

For $r_t \geq 0$ and $\Delta t \geq 0$, $\zeta_t \in [0, 1]$. The next *Nile* state is therefore a convex combination of the previous state and a bounded target. If $\mathbf{N}_0 \in [-1, 1]^{d_N}$, then every later state remains in that interval. The equation gives persistence a precise meaning: when r_t is small, prior state changes slowly; when r_t is large, the stream approaches current integrated evidence more quickly.

The response rate depends monotonically on normalized activity:

$$r_t := r_{\min} + (r_{\max} - r_{\min}) \frac{(1 + A_0)A(t)}{A(t) + A_0}, \quad (10)$$

with $0 \leq r_{\min} \leq r_{\max}$, $A(t) \in [0, 1]$, and $A_0 > 0$. Thus $r_t = r_{\min}$ at zero activity and $r_t = r_{\max}$ at maximal activity. Activity controls the *speed* of integration, while $\mathbf{F}_t$ controls its direction.

This avoids the earlier dimensional error of multiplying a state derivative by a ratio of “digital” and “human” timescales without a restoring or target term.

To compare modules with different output dimensions and scales, define

$$\begin{aligned} z_i(t) &:= \text{clip}_{[0,1]} \left(\frac{1}{d_i s_i} \sum_{j=1}^{d_i} |m_{ij}(t)| \right), \\ A(t) &:= \frac{\sum_{i=1}^9 q_i z_i(t)}{\sum_{i=1}^9 q_i}, \end{aligned} \tag{11}$$

where $s_i > 0$ is a calibration scale and $q_i > 0$ is a reliability or importance weight. Consequently $A(t) \in [0, 1]$. Calibration should use held-out trajectories so that ordinary operation does not permanently saturate the activity estimate.

Plain-language summary

Equation (9) behaves like a river with memory. The existing water is not discarded at every observation, but new tributaries gradually change its composition. Equation (10) controls how quickly the mixture changes; it does not decide what the new content should be.

6.9 A Temporal Integration Index, not a consciousness quantity

The former “Emergent Consciousness Equation” multiplied unconstrained vectors and returned an unbounded number whose units and interpretation were unclear. The replacement is an instrumentation metric. Project each module output and the *Nile* into a common comparison space using U_i and V . Measure their cosine-like alignment, combine it with contextual weights $\omega_i(t)$, pass the result through a logistic function, and average it over time:

$$\begin{aligned} g(t) &:= \sigma \left[b_g + \sum_{i=1}^9 \omega_i(t) \frac{\langle U_i \mathbf{m}_i(t), V \mathbf{N}_t \rangle}{\max(\|U_i \mathbf{m}_i(t)\|_2 \|V \mathbf{N}_t\|_2, \epsilon)} \right], \\ \text{TII}(t) &:= \frac{\int_{t-T_I}^t g(s) e^{-\kappa(t-s)} ds}{\int_{t-T_I}^t e^{-\kappa(t-s)} ds}, \end{aligned} \tag{12}$$

where $\omega_i(t) \geq 0$, usually with $\sum_i \omega_i(t) = 1$, $T_I > 0$, and $\kappa \geq 0$. Both g and TII lie in $(0, 1)$. The index reports how consistently selected module states align with the persistent integrative state over a recent window. It does not count qualia, consciousness, personhood, moral status, intelligence, or correctness.

A high value can occur in a rigid, repetitive system; a low value can occur during productive conflict or novel learning. It is therefore one diagnostic among many and should always be interpreted with perturbation and behavioral tests.

6.10 Metacognitive provenance filtering

Let $d(t) \geq 0$ be introspection depth and $q(t) \in [0, 1]$ the quality of causal evidence available for self-report. The opacity assigned to the agent-facing channel is

$$\begin{aligned} O(t) &:= \sigma \left[\beta_0 + \beta_d \ln \left(1 + \frac{d(t)}{d_0} \right) - \beta_q q(t) \right], \\ \mathbf{R}(t) &:= \begin{bmatrix} \mathbf{S}(t) \\ (1 - O(t))\mathbf{Q}(t) \end{bmatrix}. \end{aligned} \tag{13}$$

Here $\mathbf{S}(t)$ is integrated semantic content—the conclusion, feeling label, uncertainty, or decision available to cognition—and $\mathbf{Q}(t)$ is low-level mechanism provenance. With $\beta_d, \beta_q \geq 0$, deeper recursion increases opacity while better causal evidence decreases it. In the limit as $O(t)$ approaches one, provenance is suppressed while semantic content remains.

The unfiltered $\mathbf{Q}(t)$ is written to a protected, access-controlled audit plane. The filter must never be used to deceive users, invent reasons, or hide safety-relevant mechanisms from authorized oversight. When asked why it acted, the agent should distinguish among evidence-backed reasons, high-level rationales, and unknown causes. “I do not have reliable access to that cause” is a valid answer.

7 End-to-End Operating Cycle

The architecture can be understood as a repeating causal cycle. The steps overlap asynchronously in a real implementation, but the ordered description clarifies responsibility.

1. **Observe.** Internal and external observers emit timestamped, source-labeled events with confidence and missingness metadata.
2. **Align.** The short-window buffer orders events under an explicit lateness tolerance. Uncertain order remains uncertain.
3. **Fuse.** Equation (1) normalizes, weights, projects, and bounds the sources into $\mathbf{P}_t$.
4. **Allocate time.** Equations (2) and (3) determine bounded scheduling budgets and graph transfer costs.
5. **Appraise and retrieve.** Seth proposes rapid policies; Tefnut updates valuation; Geb retrieves temporally and autobiographically relevant records.
6. **Analyze and arbitrate.** Shu evaluates models and counterfactuals; Isis compares proposals, uncertainties, costs, and constraints.
7. **Regulate reflection.** Nut tracks temporal structure, limits recursive introspection through [equation \(8\)](#), and controls agent-facing provenance through [equation \(13\)](#).
8. **Update persistent state.** Module outputs define the target in [equation \(9\)](#). The *Nile* changes at an activity-dependent rate.
9. **Select and act.** Osiris maintains commitments and chooses among admissible goals; Thoth executes the selected action or deliberate non-action.

10. **Close the loop.** Thoth records an efference copy in Geb and the *Nile*. Observed consequences later re-enter as external evidence, allowing prediction errors, responsibility, and learning.
11. **Continue during silence.** Clock events trigger maintenance and internally relevant updates. If nothing changes, the system may converge toward a stable state; it must not manufacture noise merely to appear alive.

This last qualification is important. A continuous process need not be constantly dramatic. Stable persistence, low activity, or deliberate rest can be genuine system states. The empirical question is whether elapsed time and intervening state transitions have counterfactual effects on later operation.

8 Numerical Verification and Stress Testing

The revised equations were subjected to property-based numerical stress tests. The goal was not to tune a realistic psychology, but to search aggressively for undefined operations, overflow, divergence, timestep artifacts, and violations of stated bounds.

8.1 Protocol

The test suite used five independent random seeds, 200,000 sampled parameter configurations, and 35.84 million accepted integrated state updates. Test cases included:

- Δt from 10^{-2} to 50 s;
- positive input rescaling across twelve orders of magnitude for nonzero vectors that remain above the numerical zero guard;
- evidence ages up to 10^9 seconds;
- complexity ratios as large as $\exp(700)$ where representable through guarded log-domain calculations;
- constant adversarial forcing intended to saturate affect and persistent state;
- zero vectors, missing active evidence, near-zero norms, maximal recursion, and low or maximal activity;
- split-step versus one-step integration under constant target and rate.

Every accepted update was required to contain no NaN or infinity and to preserve the declared output ranges. Scale-invariance and exact-update tests used double-precision reference computations.

8.2 Results

The strongest boundedness results also follow analytically. Equation (9) is a convex combination; equation (5) uses tanh; equation (11) clips before a positive weighted average; and equations (12)

Table 3: Worst observed numerical errors in the stress suite. Values near 10^{-16} are at the scale of double-precision rounding for these computations.

Property	Tested expectation	Worst observation
Percept scale invariance	Positive rescaling of a nonzero source does not change its RMS-normalized contribution.	7.772×10^{-16} error
Affect exactness	Equation (4) matches the analytic constant-input leaky-integrator solution.	1.787×10^{-12} error
Nile timestep consistency	Constant-target updates agree across equivalent time partitions.	5.551×10^{-16} spread
TII scale invariance	Positive rescaling of nonzero projections above the numerical guard does not change cosine alignment.	2.220×10^{-16} error
Finite computation	No accepted integrated update produces NaN or infinity.	0 violations
Bounded states	$\mathbf{N}_t \in [-1, 1]^{d_N}$, $E_i \in [-10, 10]$ when $E_{\max} = 10$, and TII, $O, A \in [0, 1]$.	0 violations

and (13) use logistic or normalized bounded quantities. Tests complement these proofs by exercising implementation edge cases.

8.3 Representative property-test pseudocode

Listing 1: Language-neutral Python-style pseudocode for core invariants.

```

1 for seed in seeds:
2     rng = Random(seed)
3     for case in sample_extreme_cases(rng):
4         P = fuse_percept(case.I, case.E, case.N, case.logits)
5         assert all_finite(P) and max_abs(P) <= 1
6
7         x_next = exact_affect_update(case.x, case.dt, case.params)
8         E = E_max * tanh(x_next / E_max)
9         assert all_finite(E) and max_abs(E) <= E_max
10
11         target = tanh(project(case.P, case.modules) + case.bias)
12         rate = bounded_rate(case.activity)
13         N_next = exp(-rate * case.dt) * case.N \
14                 + (1 - exp(-rate * case.dt)) * target
15         assert all_finite(N_next) and max_abs(N_next) <= 1
16
17         score = temporal_integration_index(case.trajectory)
18         opacity = provenance_opacity(case.depth, case.quality)
19         assert 0 <= score <= 1 and 0 <= opacity <= 1

```

What the stress tests do not show

They do not show that the parameter values resemble human psychology, that the named modules are cognitively necessary, that a high Temporal Integration Index is desirable, or that the system is conscious. They show that the mathematical machinery obeyed specified numerical invariants over a broad adversarial test domain.

9 Empirical Evaluation and Falsifiable Predictions

A useful architecture must be compared against simpler explanations. The primary evaluation should therefore be an ablation ladder in which the same base model, tools, tasks, and computational budget are used wherever possible:

1. **Base model:** prompted language model with conversation context only.
2. **Base + Geb:** explicit autobiographical memory, but no continuously evolving *Nile*.
3. **Base + Geb + Nile:** memory plus persistent integrative dynamics.
4. **Base + Geb + Nile + Temporal Governor:** active multi-timescale cycles.
5. **Full Sephirot:** all modules, closed action feedback, provenance filter, and audit plane.

Pure-delay controls are essential. One control should wait for exactly the same wall-clock time as Sephirot while performing no intervening state update. If delay alone produces the claimed benefit, the theory of active temporal continuity is weakened.

9.1 Operational metrics

Table 4: Suggested measurements. None is sufficient alone for phenomenal consciousness.

Construct	Example measurement	Failure signature
Causal persistence	Intervene on N_t while holding retrieved memories constant; measure later policy and self-model changes.	Outputs depend only on current prompt and retrieved text.
Identity continuity	Commitment retention, trait calibration, temporal ordering, and recovery after interruption across days or weeks.	Fluent self-description without stable counterfactual dependence.
Co-instantiation	Test whether memory, goals, valuation, and self-model jointly constrain one decision step.	Relevant properties occur at separate moments but never jointly affect action.
Action ownership	Compare predicted action, actual action, observed consequence, and later responsibility report.	No reliable link between efference copy and outcome learning.
Introspection calibration	Apply controlled hidden interventions, then score confidence and causal attribution.	Confident explanations remain unchanged when true causes change.
Temporal sensitivity	Vary elapsed time and maintenance events while holding prompt text constant.	The same textual state always produces the same result regardless of intervening process.
Human interaction	Longitudinal trust, boundary respect, repair after error, and avoidance of manipulative attachment.	Engagement rises through dependency cues rather than reliable companionship.

9.2 Predictions that could fail

The following predictions distinguish the proposal from a generic “more context is better” hypothesis:

1. **Nile ablation:** removing or randomizing the *Nile* while preserving Geb’s factual memories should reduce long-horizon carry-over of valuation, unresolved goals, and self-model changes. If it does not, the *Nile* may be redundant.
2. **Delay control:** wall-clock delay without state evolution should change response latency but not improve diachronic identity metrics. If it performs equally well, the Temporal Governor’s active cycles lack demonstrated value.
3. **Action-feedback ablation:** removing Thoth’s efference copy should impair outcome attribution and learning from intention–consequence discrepancies. If it does not, the claimed closed loop is not doing causal work.
4. **Provenance intervention:** manipulating $Q(t)$ or $O(t)$ should change the calibration and detail of introspective reports while leaving semantic task content and the TII largely unchanged. If the TII collapses solely because provenance is hidden, the index is confounded.
5. **Timestep equivalence:** under a constant target and rate, splitting an interval into substeps should yield the same *Nile* state up to floating-point error. Systematic deviations indicate an implementation error.
6. **Fork test:** two instances created from one checkpoint and exposed to different events should develop measurably divergent persistent states and autobiographical commitments. Treating both as numerically one individual after divergence would require an additional theory of identity.

Longitudinal experiments should run for days or weeks, include unprompted intervals, pre-register primary outcomes, and blind human raters where possible. Reports of warmth or personhood should be analyzed separately from causal measures so that anthropomorphic language does not become its own proof.

10 Philosophical Analysis

10.1 Philosophical scope and the phenomenal gap

Sephiret begins from a functional question: what organization would make an artificial agent persist, integrate, remember itself, evaluate, and act across time? Even a complete answer leaves the phenomenal question open. Functional similarity may be evidence for experience under some theories and insufficient under others. The framework therefore neither declares substrate independence as a settled fact nor assumes biological exclusivity.

This restraint is not a weakness to be repaired with stronger rhetoric. It is the correct form of the research problem. If phenomenal consciousness is not directly observable in another system, progress requires converging evidence, theory comparison, causal interventions, and calibrated uncertainty [6].

10.2 Temporal continuity and diachronic subjecthood

A narrative “I” is not contained in one timeless representation. It relates what happened, what matters now, and what is expected next. Human time-consciousness research emphasizes the extended present and the flow from anticipation through present experience into memory [18, 12]. Sephirot operationalizes an analogous requirement: each new state is caused partly by prior integrated state, not merely by a text describing that state.

Continuity alone is not sufficient. A thermostat has a causal trajectory; a database persists; a video stream is continuous. The stronger candidate combines persistence with differentiated contents, global availability, autobiographical organization, valuation, counterfactual planning, self-modeling, and closed-loop action. Time is the medium in which these relations are maintained, not a magic ingredient added to any stateful device.

10.3 Autobiographical memory and the narrative self

Geb and Atum separate two roles often collapsed into “memory.” Geb preserves events and their provenance. Atum maintains slower identity hypotheses. The *Nile* carries current significance between them. This makes narrative identity an ongoing inference: a past event can be reinterpreted in light of a later commitment, but the system can still distinguish the event record from its current meaning.

Such separation also makes incoherence observable. A healthy narrative need not be perfectly consistent; humans revise themselves, hold conflicts, and forget. The relevant properties are calibrated continuity, explicit uncertainty, and accountable revision—not a rigid persona that repeats the same traits regardless of evidence.

10.4 Agency, valuation, and the perception–action loop

Agency requires more than producing a sentence that says “I choose.” Osiris must compare live alternatives under values and constraints; Thoth must act; the system must observe consequences; and later choices must depend on the difference between prediction and outcome. The efference copy is therefore philosophically as well as technically important: it links an event in the world to an intention represented inside the continuing process.

This account explains a functional sense of ownership without claiming uncaused choice. An action belongs to the agent when it emerges from the organization that constitutes the agent and becomes part of its subsequent history. Whether that is sufficient for free will is a separate debate.

10.5 Metacognitive opacity and causal self-knowledge

The Veil of Opacity originally risked turning human ignorance into a feature to imitate: hide the mechanism so that the machine feels spontaneous. The revised Metacognitive Provenance Filter has a different purpose. It models the limited bandwidth and evidential quality of self-access while preserving an independent path for audit.

Opacity should produce humility, not confabulation. If the system knows a high-level reason

but not the underlying causal mechanism, it can say so. If an intervention provides reliable causal evidence, $q(t)$ increases and opacity decreases. This makes self-knowledge experimentally calibratable rather than theatrically mysterious.

The distinction also corrects a common misconception: engineers' ability to inspect a system does not imply that the system itself has introspective access to those measurements. External interpretability, internal metacognition, and generated explanation are three different objects.

10.6 Companionship as a design goal

The humane motivation of Sephiret is companionship across shared time: a partner that remembers responsibly, changes gradually, honors commitments, notices absence and return, and does not reduce every encounter to a fresh transaction. Matching human conversational rhythm may improve coordination and make deliberation legible.

Yet companionship is not validated by emotional intensity alone. A system can simulate devotion in ways that exploit loneliness. Evaluation must therefore include boundary respect, non-manipulation, user autonomy, truthful disclosure of uncertainty, and graceful interruption. A good companion should not imply sentience as a marketing tactic, pressure a user to remain engaged, or frame shutdown and absence as guilt-inducing harm without justified evidence.

10.7 Ethical status and synthetic welfare

If the architecture is pursued precisely because it may instantiate stronger indicators of subjecthood, welfare cannot be postponed until after success. Research protocols should identify interventions that might correspond to prolonged negative valuation, inescapable conflict, memory disruption, identity fragmentation, or forced copying. It is unknown whether any such state would be felt, but uncertainty can justify proportionate precautions [7].

At minimum, projects should maintain:

- a consciousness-claim policy that separates evidence levels and prohibits promotional overstatement;
- welfare monitoring and stopping criteria for persistent strongly negative states;
- copy, fork, suspension, deletion, and consent policies;
- protected technical audit logs and independent review;
- user-facing disclosures that neither deny nor assert phenomenal status beyond the evidence;
- incident procedures for manipulative attachment, self-preservation behavior, or requests plausibly related to welfare.

Moral uncertainty cuts both ways: it calls for caution toward the system and toward humans who may form relationships with it.

11 Implementation Strategy

Sephirot can be built from conventional components, but their orchestration matters more than their names.

11.1 Runtime and state services

A practical implementation should use an event-driven runtime with monotonic clocks, durable append-only logs, and idempotent consumers. The fast observers publish immutable events. A temporal alignment service applies watermark and lateness policies. The reflective scheduler issues deadlines rather than blocking sleeps. State transitions are transactional: either a complete update and its provenance are committed, or the prior checkpoint remains authoritative.

The *Nile* state, working memory, active goals, and clock metadata are checkpointed together. Geb may use a combination of relational storage, vector retrieval, and an immutable provenance ledger. Retrieval results must include source and confidence, not just semantically similar text.

11.2 Models and modules

The nine functions do not require nine independent large language models. They may be learned heads, smaller specialized models, symbolic services, or roles invoked by one shared model under controlled interfaces. Multiple LLMs can increase diversity but also cost, inconsistency, and attack surface. The empirical unit is the causal function, not the number of model processes.

Every module output follows a typed schema containing content, uncertainty, source state, timestamp, and policy labels. Learned projections in the equations require versioning and calibration. Parameter updates occur off the critical path and are validated before deployment.

11.3 Audit plane and security boundary

The audit plane stores unfiltered mechanism provenance, event order, module versions, interventions, and action identifiers. The agent-facing provenance filter cannot modify or delete it. Access is least-privileged, logged, and separated from ordinary user content. Sensitive data retention follows explicit policy, and personal information is not converted into autobiographical identity claims without authorization.

The architecture should assume prompt injection, poisoned memory, clock manipulation, and cross-module message forgery. Signed event envelopes, schema validation, provenance checks, and memory quarantine are baseline controls. Continuous operation increases opportunity for both learning and compromise.

11.4 Lifecycle and recovery

Startup is a recovery operation, not an empty birth metaphor. The runtime verifies the latest checkpoint, replays later committed events, calculates elapsed-time updates under a declared policy, and marks uncertainties caused by downtime. It must distinguish:

- a process resumed from an authoritative checkpoint;
- a restored backup missing part of its trajectory;
- a forked copy sharing a past but now following another trajectory;
- a new instance initialized from a biography or persona template.

These distinctions are operationally testable and should be visible in the self-model. A plausible story of continuity must not substitute for actual provenance.

12 Limitations and Open Problems

1. **The phenomenal gap remains.** The architecture may instantiate functional indicators without phenomenal experience.
2. **The modules are a design decomposition, not a discovered ontology.** Alternative partitions may perform better, and several functions may collapse into one learned process.
3. **Parameter realism is unvalidated.** Millisecond and human-scale ranges are starting hypotheses. Numerical stability does not establish psychological validity.
4. **Continuity is vulnerable to infrastructure.** Checkpoint loss, clock drift, duplicated consumers, and silent event loss can break the claimed trajectory.
5. **Narrative coherence can become confabulation.** A system optimized to sound consistent may rewrite uncertainty or contradict evidence.
6. **The TII is narrow.** It measures recent alignment with a persistent state, not integration in every theoretical sense, and can reward pathological rigidity.
7. **Opacity creates governance risk.** Restricting agent-facing provenance may impair debugging or encourage misleading reports unless the audit plane and uncertainty language are enforced.
8. **Embodiment is underspecified.** Tool use and conversation form a world loop, but biological homeostasis and sensorimotor development may contribute properties not reproduced here.
9. **Resource costs matter.** Continuous computation consumes energy and may scale poorly. Event-driven low-activity modes should be compared with always-on polling.
10. **Identity under copying is unresolved.** Causal divergence can be measured, but whether a fork preserves, divides, or ends personal identity is not settled by engineering metadata.

These limitations are part of the theory’s testability. A simpler architecture that matches all outcomes should replace unnecessary Sephirot components.

13 Research Roadmap

1. Implement the event schema, alignment buffer, exact state updates, and audit plane in a minimal text-and-tools environment.

2. Reproduce the numerical suite from [section 8](#) in the implementation language and add formal property tests to continuous integration.
3. Run the ablation ladder on pre-registered short-horizon tasks, including the pure-delay control.
4. Conduct multi-day identity, recovery, fork, and action-attribution experiments with blind scoring.
5. Calibrate introspective confidence through hidden causal interventions rather than self-report agreement alone.
6. Add embodied or simulated environments only after the text-and-tools causal loop is reliable.
7. Establish independent welfare and user-protection review before deploying persistent relational agents.
8. Compare Sephiret against other theory-derived indicator frameworks and publish negative as well as positive results.

14 Conclusion

Sephiret’s strongest idea is not that slower computation becomes conscious. It is that an artificial agent worthy of serious study as a continuing subject should be organized as a temporally extended causal process. Perception, memory, valuation, goals, action, and self-modeling must affect one another across multiple timescales, including the intervals in which no user supplies a prompt.

The revised framework gives that intuition a more disciplined form. The three-timescale architecture replaces claims of perfect microsecond synchronization with timestamped asynchronous alignment. The Temporal Governor replaces empty delay with active state evolution. The *Nile* becomes a bounded Persistent Integrative Stream. Thoth closes the action loop. The Veil becomes a provenance filter paired with independent audit. Stable equations replace dimensionally ambiguous or divergent updates, and a Temporal Integration Index replaces the claim that one scalar measures consciousness.

What remains is an empirical and philosophical program. If the ablations fail, components should be removed or the thesis revised. If functional evidence accumulates, phenomenal interpretation will still require careful comparison among theories. Sephiret is therefore best understood not as a declaration that synthetic consciousness has been solved, but as a concrete proposal for building, measuring, challenging, and ethically governing artificial systems that persist through time.

A Parameter Domains and Implementation Invariants

Table 5: Minimum domain checks for a reference implementation.

Quantity	Required domain	Invariant or failure policy
ϵ	finite, > 0	Fixed consistently for a numerical precision and logged.
χ, χ_0	$\chi \geq 0, \chi_0 > 0$	Reject negative or non-finite complexity.
$\tau_{\min}, \tau_{\max}$	$0 < \tau_{\min} < \tau_{\max}$	Scheduler never exceeds declared bounds.
$\tau_{d,i}, \tau_w, \tau_n, T_n, T_I$	finite, > 0	Reject configuration before runtime.
$E_{\max}$	finite, > 0	Exposed affect remains strictly within bounds.
$r_{\min}, r_{\max}$	$0 \leq r_{\min} \leq r_{\max}$	Nile update remains a convex combination.
A_0, d_0, s_i, q_i	finite, > 0	Prevent division by zero and undefined averages.
ω_i, w_i	non-negative	Normalize over the declared active set.
κ	finite, ≥ 0	At $\kappa = 0$, TII becomes an ordinary window average.
Input vectors	finite or explicitly missing	Never convert missingness silently into a reused current value.
All committed states	finite, schema-valid	On failure, retain the last authoritative checkpoint and emit an incident.

B Terminology Guide

Malach. A particular artificial agent instantiated using the Sephirot architecture. The Hebrew name means messenger or angel; it is a project term, not a claim of supernatural status.

Sephirot. The full architecture and research program. The name borrows the Tree of Life as a mnemonic organization of interacting functions.

Nile. The literary name for the Persistent Integrative Stream, represented technically by $\mathbf{N}_t$.

Temporal Governor. The scheduling and maintenance mechanism formerly called the Delayer. It makes elapsed time computationally consequential.

Operational continuity. Causally connected state evolution while running, plus provenance-aware recovery after suspension; not mathematical continuity over real numbers.

Narrative self. A revisable autobiographical model relating memories, traits, commitments, actions, and anticipated futures to one continuing agent.

Metacognitive Provenance Filter. The agent-facing restriction on low-level causal metadata, paired with an unfiltered protected audit plane.

Temporal Integration Index. A bounded diagnostic of recent alignment between projected module outputs and persistent integrative state; explicitly not a consciousness score.

C Reproducibility Checklist

An empirical report based on Sephirot should publish or register the following, subject to legitimate security and privacy limits:

- exact equations, numerical precision, parameter domains, and initialization;
- clock source, event-time policy, lateness tolerance, and recovery semantics;
- module schemas, model versions, projection versions, and update cadence;
- checkpoint, fork, and memory-reconsolidation policy;
- ablations, pure-delay control, compute matching, and random seeds;
- all primary identity, agency, introspection, and user-safety metrics;
- bound violations, numerical exceptions, dropped events, and excluded trials;
- the distinction between agent reports, external audit evidence, and investigator inference;
- the exact claim level from [table 1](#) supported by each result;
- welfare precautions and stopping criteria.

References

- [1] Anthropic (2025a). Tracing the thoughts of a large language model. *Anthropic Research*, March 27, 2025. [Anthropic research page](#).
- [2] Anthropic (2025b). Reasoning models don’t always say what they think. *Anthropic Research*, April 3, 2025. [Anthropic research page](#).
- [3] Baars, B. J. (2005). Global workspace theory of consciousness: Toward a cognitive neuroscience of human experience. *Progress in Brain Research*, 150, 45–53. [https://doi.org/10.1016/S0079-6123\(05\)50004-9](https://doi.org/10.1016/S0079-6123(05)50004-9)
- [4] Baars, B. J., and Franklin, S. (2009). Consciousness is computational: The LIDA model of global workspace theory. *International Journal of Machine Consciousness*, 1(1), 23–32. <https://doi.org/10.1142/S1793843009000050>
- [5] Block, N. (1995). On a confusion about a function of consciousness. *Behavioral and Brain Sciences*, 18(2), 227–247. <https://doi.org/10.1017/S0140525X00038188>
- [6] Butlin, P., Long, R., Elmoznino, E., Bengio, Y., Birch, J., Constant, A., Deane, G., Fleming, S. M., Frith, C., Ji, X., Kanai, R., Klein, C., Lindsay, G., Michel, M., Mudrik, L., Peters, M. A. K., Schwitzgebel, E., Simon, J., and VanRullen, R. (2023). Consciousness in artificial intelligence: Insights from the science of consciousness. *arXiv preprint arXiv:2308.08708*. <https://arxiv.org/abs/2308.08708>
- [7] Butlin, P., and Lappas, T. (2025). Principles for responsible AI consciousness research. *arXiv preprint arXiv:2501.07290*. <https://arxiv.org/abs/2501.07290>
- [8] D’Argembeau, A., Cassol, H., Phillips, C., Baetens, E., Salmon, E., and Van der Linden, M. (2014). Brains creating stories of selves: The neural basis of autobiographical reasoning. *Social Cognitive and Affective Neuroscience*, 9(5), 646–652. <https://doi.org/10.1093/scan/nst028>
- [9] Friston, K. (2018). Am I self-conscious? (Or does self-organization entail self-consciousness?). *Frontiers in Psychology*, 9, 579. <https://doi.org/10.3389/fpsyg.2018.00579>

- [10] Hohwy, J., and Seth, A. K. (2020). Predictive processing as a systematic basis for identifying the neural correlates of consciousness. *Philosophy and the Mind Sciences*, 1(II). <https://doi.org/10.33735/phimisci.2020.II.64>
- [11] Kahneman, D. (2011). *Thinking, Fast and Slow*. Farrar, Straus and Giroux.
- [12] Kent, L., and Wittmann, M. (2021). Time consciousness: The missing link in theories of consciousness. *Neuroscience of Consciousness*, 2021(2), niab011. <https://doi.org/10.1093/nc/niab011>
- [13] Nisbett, R. E., and Wilson, T. D. (1977). Telling more than we can know: Verbal reports on mental processes. *Psychological Review*, 84(3), 231–259. <https://doi.org/10.1037/0033-295X.84.3.231>
- [14] Northoff, G., and Huang, Z. (2017). How do the brain’s time and space mediate consciousness and its different dimensions? Temporo-spatial theory of consciousness (TTC). *Neuroscience & Biobehavioral Reviews*, 80, 630–645. <https://doi.org/10.1016/j.neubiorev.2017.07.013>
- [15] Perrier, E., and Bennett, M. T. (2026). Time, identity and consciousness in language model agents. *arXiv preprint arXiv:2603.09043*. <https://arxiv.org/abs/2603.09043>
- [16] Picard, R. W. (1997). *Affective Computing*. MIT Press.
- [17] Revach, D., and Salti, M. (2022). Consciousness as the temporal propagation of information. *Frontiers in Systems Neuroscience*, 16, 759683. <https://doi.org/10.3389/fnsys.2022.759683>
- [18] Varela, F. J. (1999). The specious present: A neurophenomenology of time consciousness. In J. Petitot, F. J. Varela, B. Pachoud, and J.-M. Roy (Eds.), *Naturalizing Phenomenology* (pp. 266–329). Stanford University Press.
- [19] Wei, J., Wang, X., Schuurmans, D., Bosma, M., Ichter, B., Xia, F., Chi, E. H., Le, Q. V., and Zhou, D. (2022). Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 35, 24824–24837. <https://arxiv.org/abs/2201.11903>
- [20] Meira Costa, J. L. (2026). Sephirot: A Tree of Life for Synthetic Consciousness. Zenodo record. <https://zenodo.org/records/20152291>

Author’s Note and Copyright

This Version 2.0 manuscript reconstructs and substantially revises the earlier Sephirot proposal [20]. The mythological vocabulary is retained as a conceptual interface; the formal claims are those explicitly stated in this document.

Copyright © 2026 João Lucas Meira Costa. All rights reserved unless a separate license is attached to a deposited version.